

Lesson: Module 1 - Establishing Your Baseline

Video: M1V3 – Demo: Running Your First Baseline Benchmark

Purpose: Run a minimal benchmarking harness to capture TTFT, latency, throughput, and cost metrics against a fixed evaluation set.**Commands:**

```
# Navigate to project

cd baseline_demo

# Install dependencies

pip install -r requirements.txt

# Set OpenAI API key

export OPENAI_API_KEY="your-key-here"

# Run the baseline benchmark

python run_baseline.py --model gpt-4o-mini --output baseline_results
```

Key Files to Open:

- eval_set.jsonl – Fixed evaluation prompts (5 tasks)
- run_baseline.py – Benchmarking harness with streaming + timing
- baseline_results/raw_logs.jsonl – Detailed per-request logs
- baseline_results/summary.csv – Clean metrics for analysis
- notes.txt – Outlier observations

Key Metrics Captured:

- TTFT (time-to-first-token): 0.4s – 1.1s
- Total latency: 0.5s – 4.4s
- Tokens per second: 12 – 42
- Cost per task: \$0.000003 – \$0.000075