

## **Lesson: Module 1 – Establishing Your Baseline**

### **Video: M1V3 – Demo: Running Your First Baseline Benchmark**

#### **A. Repository file:**

- demo1\_M1V3.zip

#### **Contents:**

- eval\_set.jsonl – Fixed evaluation prompts (5 tasks)
- run\_baseline.py – Benchmarking harness with streaming + timing
- requirements.txt – Python dependencies (openai)
- notes.txt – Template for outlier observations

#### **After running, generates:**

- baseline\_results/raw\_logs.jsonl – Detailed per-request logs
- baseline\_results/summary.csv – Clean metrics table

#### **Purpose of the Repository/Resource:**

- To provide learners with a minimal, production-style benchmarking harness demonstrated in the video.
- To enable learners to run baseline measurements locally and understand how to capture TTFT, latency, throughput, and cost metrics.
- To give learners a repeatable template they can adapt for their own LLM applications and model comparisons.