

## **Lesson: Module 3 – Run an A/B/C Test & Pick a Winner**

### **Video: M3V3 – Demo: Running Your First A/B/C BenchmarkA. Repository file:**

- demo2\_M3V3.zip

#### **Contents:**

- eval\_set.jsonl – Fixed evaluation prompts (10 tasks)
- run\_variant.py – Benchmarking harness with streaming + timing
- run\_all.sh – Optional script to run all 3 variants
- analyze\_results.ipynb – Pandas + matplotlib comparison notebook

#### **After running, generates:**

- results\_A/raw.jsonl – Detailed per-request logs
- results\_A/summary.csv – Clean metrics table
- results\_B/ – Same structure for Variant B
- results\_C/ – Same structure for Variant C

#### **Purpose of the Repository/Resource:**

- To provide learners with a clean A/B/C testing harness demonstrated in the video.
- To enable learners to run multiple variants on a fixed dataset and compare TTFT, latency, throughput, cost, and quality metrics.
- To teach learners how to visualize trade-offs using pandas and matplotlib for data-driven ship/no-ship decisions.
- To give learners a repeatable template they can adapt for their own model comparisons, prompt variants, or configuration experiments.