

Non-Graded Project: Design & Present a Responsible AI Solution Pack

Description

In this capstone project, you step into the role of a Responsible AI consultant for an online recruitment platform preparing to launch an AI-driven candidate screening model. Executives depend on this system to shortlist applicants quickly, but are worried about fairness, bias, and explainability. You will use structured Responsible AI tools to map stakeholders and risks, identify potential bias points across the AI lifecycle, design a targeted Bias Mitigation Plan, and create a stakeholder-ready presentation that clearly communicates your ethical design decisions.

Learning Objectives

By the end of this project, learners will be able to:

- Map stakeholders, risks, and mitigations for an AI-driven recruitment system using a Responsible AI Canvas.
- Identify and document potential bias points across data collection, model training, and decision stages.
- Develop a concise Bias Mitigation Plan that addresses data-level, model-level, and process-level interventions.



- Create a stakeholder-ready pitch deck that explains key risks, safeguards, and governance outcomes in clear, non-technical language.
- Synthesize Responsible AI design, evaluation, and communication practices into a coherent solution pack for a real-world scenario.



Project Components

Introductory Slide

Title:

Hiring at Scale, Ethically: Responsible AI for Candidate Screening

Subtitle:

Design a fair, explainable screening model and present it with stakeholder-ready clarity

Tools to be used:

→ **Excalidraw** - for designing a Responsible AI Canvas (stakeholders, risks, mitigations, data flows).

→ **Gamma.app** - for generating a 5-6 slide stakeholder presentation.

→ Optional support tools: any text editor (for notes), diagram export tools (PNG/PDF export from Excalidraw and Gamma).

Scenario

TalentFlow, an online recruitment platform, is about to launch an AI-driven candidate screening model that automatically evaluates resumes and produces ranked shortlists for recruiters. Executives want faster hiring cycles and reduced manual screening time, but concerns have emerged:

- Will the model disadvantage candidates from certain schools, regions, or backgrounds?
- Can recruiters and candidates understand why someone was shortlisted or rejected?

- What happens if bias is discovered after deployment?

You have been hired as the **Responsible AI consultant** to design and present a Responsible AI Solution Pack. Your work will inform the final go/no-go decision for launch, and will be presented to executives, legal/compliance, and HR leadership. They are not interested in raw model details alone—they want evidence that fairness, transparency, and accountability have been built into the system from the start.

Instructions for Learners

Step 1: Map the Responsible AI Canvas in Excalidraw

1

Open Excalidraw and create a blank canvas.

2

Sketch a Responsible AI Canvas for the TalentFlow screening model, including:

Key stakeholders (recruiters, candidates, HR, legal, affected communities, regulators).

Main data sources (resumes, applicant tracking system data, historical hiring decisions, optional assessments).



Model decisions (ranking, shortlisting, automated rejections, recruiter-facing explanations).

Potential harms (unfair ranking by gender or school, opaque rejections, feedback loops reinforcing past bias).

3

For each high-risk area, note at least one mitigation (e.g., fairness audits, explanation UI, appeals process).

4

Clearly label each section so another person could understand your canvas at a glance.

Tip: Think of the canvas as a one-page X-ray of the system: who it touches, where risk lives, and how you plan to manage it.

Step 2: Identify Bias Points Across the AI Lifecycle

1 Using your canvas as a guide, list **3-5 potential bias points** across:

→ Data collection (e.g., historic hiring skew toward certain schools or demographics).

→ Model training and feature selection (e.g., using proxies like school name or zip code).

→ Decision and post-decision stages (e.g., automated rejection emails with no appeal path).

- 2 For each bias point, specify which Responsible AI principles are most at risk (Fairness, Transparency, Accountability, Privacy, Safety).
- 3 Keep this list as the foundation for your Bias Mitigation Plan in the next step.

Timebox: Aim to spend no more than 10 minutes on this analysis so you have time for planning and presentation.

Step 3: Create a Bias Mitigation Plan

1 Draft a **1-2 page Bias Mitigation Plan** that includes:

→ **Data-level interventions** (e.g., rebalancing training data, excluding sensitive proxies, running pre-training fairness checks).

→ **Model-level interventions** (e.g., fairness constraints, subgroup performance monitoring, explanation techniques).

→ **Process-level interventions** (e.g., human-in-the-loop review for borderline cases, appeals process, periodic audits).

2 For each mitigation, briefly describe:

→ What it is

→ Which risk it addresses

→ How success would be measured (e.g., reduced disparity between groups, improved explanation quality).

3 Ensure the language is clear enough that HR or legal stakeholders could read it without technical background.

Tip: Think like a consultant: concise, specific, and tied to observable outcomes.

Step 4: Build a 5-6 Slide Stakeholder Presentation in Gamma.app

- 1 Open Gamma.app and create a new presentation.
- 2 Design 5-6 concise slides that a busy executive could understand in under 5 minutes. Suggested structure:
 - Slide 1: Problem & Purpose - why TalentFlow needs an AI screening model and who it affects.
 - Slide 2: Key Risks - fairness, transparency, and reputational risks you identified.



Slide 3: Bias Mitigation Plan - top interventions across data, model, and process.

Slide 4: Governance & Oversight - who is accountable, audit processes, documentation (Model Cards, logs, appeals).

Slide 5: Expected Outcomes - efficiency improvements plus fairness and transparency commitments.

Optional Slide 6: Open Questions & Next Steps - remaining risks, pilot plans, additional monitoring.

3 Use short sentences, clean visuals, and avoid jargon.

4 Export your slides as **PDF** or **image files** for submission.

Tip: Imagine presenting this to a skeptical board. Every slide should answer the question: “Why should we trust this system?”

Step 5: Reflect on Your Responsible AI Solution Pack

1 Write a short reflection (200-300 words) explaining:

→ Which part of the work (canvas, mitigation plan, or deck) you found most challenging and why.

→ How your thinking about Responsible AI in recruitment has changed.

→ One concrete practice you would apply in a real AI project after this course.

Tip: Honest reflection is more valuable than perfection—focus on what you actually learned.

Compile Your Findings

All deliverables should be based on the TalentFlow recruitment scenario and align with the tools and steps above.

Requirement 1: Responsible AI Canvas Artifact: Responsible AI Canvas showing stakeholders, risks, data flows, and mitigations.

Include:

- Key stakeholders clearly labeled
- Main data sources and flows

- Model decisions and outputs
- Potential harms identified
- Mitigations for each high-risk area

Requirement 2: Bias Points & Mitigation Plan: Bias Mitigation Plan describing identified bias points and proposed interventions.

Include:

- 3-5 identified bias points across the AI lifecycle
- Responsible AI principles at risk for each bias point
- Data-level, model-level, and process-level interventions

- Success metrics for each mitigation
- Clear, non-technical language suitable for stakeholders

Requirement 3: Stakeholder Presentation: Stakeholder deck outlining purpose, risks, mitigations, and governance.

Include:

- Problem & Purpose slide
- Key Risks slide
- Bias Mitigation Plan slide
- Governance & Oversight slide

- Expected Outcomes slide
- Clean visuals and jargon-free language

Requirement 4: Reflection on Responsible AI Design & Communication:
Reflection on your approach and key takeaways.

Include:

- Which part of the work did you find most challenging and why
- How your thinking about Responsible AI in recruitment has changed
- One concrete practice you would apply in a real AI project
- Key learnings applicable to future assessments
- Professional growth areas identified

Summary

This course-end activity provides hands-on experience with Responsible AI design, bias mitigation, and ethical communication. Learners build competencies in stakeholder mapping, risk identification, mitigation planning, and presenting ethical safeguards in clear, accessible language—skills applicable to real-world AI governance and deployment.

Share The Project

Document: Save the report in PDF format.

Upload: Share in the “Peer Review” area of the course shell with a brief description.

Instructions for Documenting and Sharing The Project for Peer Review

1. Document the Project

- Use word processing software to create your report.
- Ensure all sections of the project document structure are completed.
- Proofread and edit your report for clarity and accuracy.

2. Share the Project

- Save your report in PDF format.
- Upload your documents to the “Peer Review” area in the course shell.
- Provide a brief description of your project in the submission post.

Instructions for Documenting and Sharing The Project for Peer Review

3. Peer Review

- Review the projects submitted by your peers.
- Provide constructive feedback and comments on their work.